

A brief intro to Machine Learning

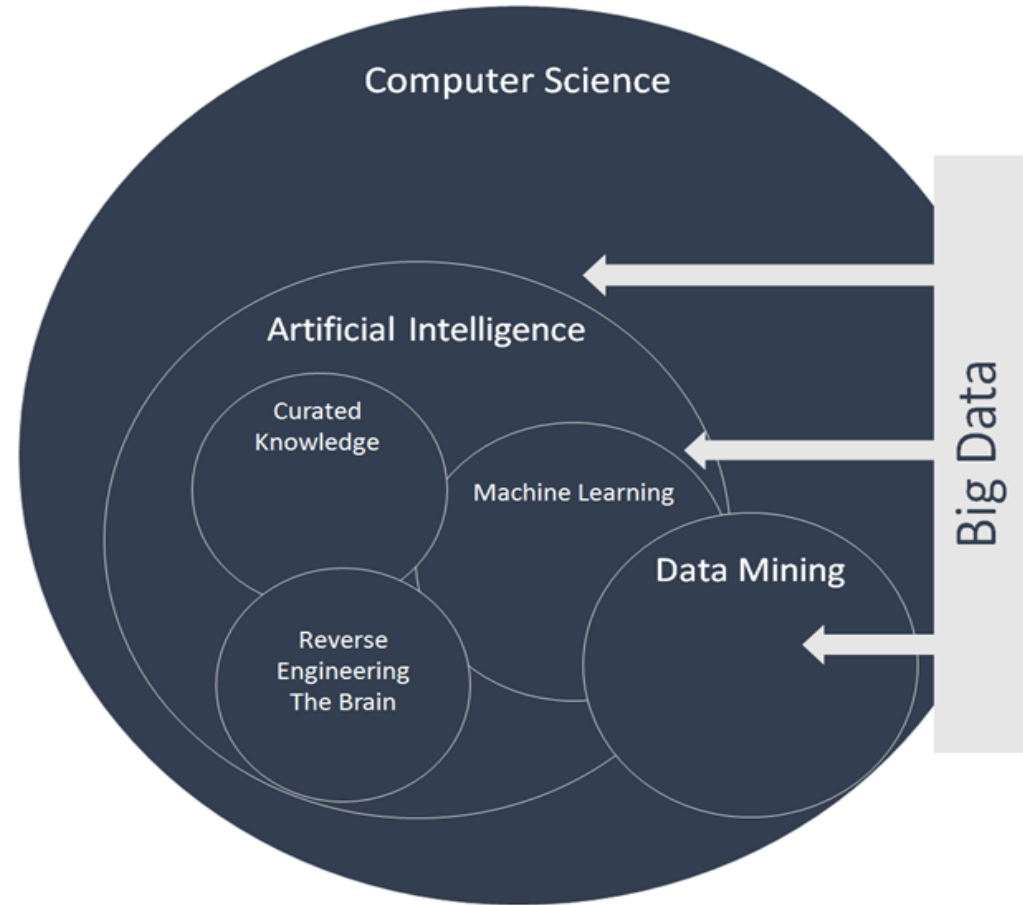
Daniel Eftekhari
daniel.eftekhari@ibm.com

Today

- Essential ML concepts
- Introduction to two ML algorithms
- How to evaluate ML algorithms

What is Machine Learning

- “... *algorithms that can learn from and make predictions on data ...*”
(Wikipedia)
- Statistical pattern recognition



[1] <https://inovancetech.com/buzzwords.html>

[1]

Machine Learning Subcategories

- Supervised Learning
 - Learn relation between data and known labels/outcomes
 - $X \rightarrow Y$
- Unsupervised Learning
 - Learn patterns in data – no labels/outcomes given
- Reinforcement Learning
 - Learn what actions maximize reward
 - Ex: robot-training, understanding biological decision-making

Today's focus – Supervised Learning

- Brief introduction to
 - Linear regression
 - Logistic regression
 - Validating results
- Knowledge of the above is good preparation for future encounters with supervised learning


Before we start

- Some general advice:
 - Occam's razor: *"Among competing hypotheses, the one with the fewest assumptions should be selected."* (Wikipedia)
 - Be wary: very easy to treat ML algorithms as black boxes
 - Particularly true as there are libraries dedicated to making life easy in many programming languages
 - Invariably leads to sub-optimal algorithm selection and/or design
 - Algorithm of choice is almost always problem-specific
 - Before applying ML algorithms
 - Understand the nature of your data -> how was it obtained?
 - Visualize your data -> is there a class bias in your data?
 - Develop baseline models -> will serve as a metric when using more advanced algorithms

Linear Regression

Linear Regression - Essentials

- Used for prediction
 - Maps input to continuous outputs
 - ex: house prices, customer ratings
 - $y \rightarrow$ Known label
- $\hat{y} = X\beta + \epsilon$
 - $\hat{y}_i \rightarrow$ *dependent variable*
 - $x_{i1} \dots x_{im} \rightarrow$ *independent variables*
 - $\beta \rightarrow$ *parameter vector*
 - $\epsilon_i \rightarrow$ *noise*




$$\hat{y}_i = \beta_0 + \beta_1 x_{i,1} + \beta_2 x_{i,2} + \dots + \beta_m x_{i,m} + \epsilon_i$$

[1] https://en.wikipedia.org/wiki/Linear_regression

Linear Regression - Essentials

- Key Assumptions of Linear Regression
 - Linearity (not in the way you think)
 - Linear in β
 - $\hat{y}_i = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \dots + \beta_m x_i^m + \epsilon_i$
 - Constant variance in error (homoscedasticity)
 - Often a flawed assumption
 - Linear independence of predictors



[1]

[1] https://en.wikipedia.org/wiki/Linear_regression

Linear Regression - Optimization

- Loss Function
 - Metric for evaluating how well model fits data
- Ordinary Least Squares
 - $loss = \frac{1}{2} (\mathbf{y} - \hat{\mathbf{y}})^2$

Linear Regression – Analytical Solution

- Minimize loss w.r.t. β
 - $loss = \frac{1}{2}(\mathbf{y} - \hat{\mathbf{y}})^2$
 - $loss = \frac{1}{2}(\mathbf{y} - \mathbf{X}\beta)^2$
 - $\frac{\partial loss}{\partial \beta} = \frac{\partial}{\partial \beta} (\frac{1}{2}(\mathbf{y} - \mathbf{X}\beta)^2)$
 - $0 = -\mathbf{X}^T(\mathbf{y} - \mathbf{X}\beta)$
 - $\mathbf{X}^T\mathbf{y} = \mathbf{X}^T\mathbf{X}\beta$
 - $\beta = (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{y}$
- Problem:
 - Computational complexity: $O(M^2N)$
 - M = number of features, N = number of samples

Linear Regression – Gradient Descent


- Algorithm

1. Initialize β randomly
2. Repeat until convergence

$$\beta \leftarrow \beta - \lambda \frac{\partial \text{loss}}{\partial \beta}$$

- λ = learning rate ($0 < \lambda < 1$)


- Too low:
 - slow to converge
 - trapped in local minima
- Too high
 - divergence




[1]

[1] https://en.wikipedia.org/wiki/Gradient_descent

Linear Regression – Gradient Descent



Large learning rate: Overshooting.



Small learning rate: Many iterations until convergence and trapping in local minima.

Can use momentum to overcome.

[1]

[1] http://sebastianraschka.com/Articles/2015_singlelayer_neurons.html



Linear Regression – Gradient Descent

- Batch updates:
 - $\boldsymbol{\beta} \leftarrow \boldsymbol{\beta} + \lambda \sum_{n=1}^N [(y^{(n)} - \hat{y}(x^{(n)}))x^{(n)}]$
- Stochastic Gradient Descent
 - Computationally faster (do not need to hold entire dataset in memory)
 - Update weights for each training case
 - *for* $i = 1$ *to* N
 - $\boldsymbol{\beta} \leftarrow \boldsymbol{\beta} + \lambda (y^{(n)} - \hat{y}(x^{(n)})) x^{(n)}$
 - Can use mini-batches: balance between performance and speed

Linear Regression – Under/Overfitting



- Underfitting

- Model does not capture complexity of problem



- Overfitting

- Model is too powerful – does not generalize to unseen data



[1] Zemel, Urtasun, Fidler (UofT), CSC411 Fall 2016, taken from Bishop, Pattern Recognition and Machine Learning, 2006

Linear Regression – Regularization

- Problem
 - As we increase model power, weights increase in magnitude to compensate for noise
- Solution: Regularization
 - Restrict magnitude of $\boldsymbol{\beta}$ by penalizing large weights
- Ridge regression
 - $loss = \frac{1}{2} (\mathbf{X}\boldsymbol{\beta} - \mathbf{y})^2 + \frac{\alpha}{2} \boldsymbol{\beta}^T \boldsymbol{\beta}$
 - α = regularization term ($0 < \alpha < 1$)

Linear Regression – Regularization

- Analytical solution:

- $\boldsymbol{\beta} = (\mathbf{X}^T \mathbf{X} + \alpha \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$

- Gradient descent:

- $\boldsymbol{\beta} \leftarrow \boldsymbol{\beta} + \lambda \sum_{n=1}^N [(y^{(n)} - \hat{y}(x^{(n)}))x^{(n)} - \alpha \boldsymbol{\beta}]$

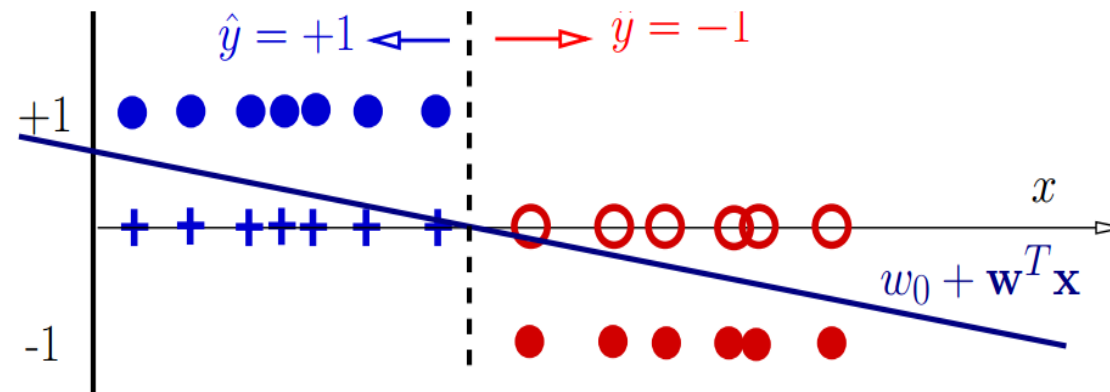
Linear Regression – Key Concepts

- Used for prediction
- Model assumptions
- Loss function
- Analytical solution vs. gradient descent
- Under/overfitting
- Hyperparameter tuning

Logistic Regression

Logistic Regression - Essentials



- Used for classification
 - Maps input to binary output
 - ex: medical diagnostics
 - We will focus on binary classification
- Linear regression (blue line) isn't a good solution
 - The sign function (dotted line), is more appropriate



Logistic Regression – Sigmoid Function

- The sign function is an extreme case of the sigmoid function
 - Sigmoid properties:
 - probabilistic outputs
 - continuous and differentiable everywhere
 - maintains classification property

- $\sigma(z) = \frac{1}{1+\exp(-z)}$




[1] https://en.wikipedia.org/wiki/Sigmoid_function

[2] <http://www.gnuplotting.org/defining-piecewise-functions/>

Logistic Regression – Binary Classification

- $\hat{y}(\mathbf{x}) = \sigma(\boldsymbol{\beta}^T \mathbf{x} + \beta_0)$
 - $p(C = 0|\mathbf{x}) = \sigma(\boldsymbol{\beta}^T \mathbf{x} + \beta_0)$
 - $p(C = 1|\mathbf{x}) = 1 - \sigma(\boldsymbol{\beta}^T \mathbf{x} + \beta_0)$
- Decision boundary:
 - $\boldsymbol{\beta}^T \mathbf{x} + \beta_0 = 0$
- Linear Regression
 - sensitive to outliers



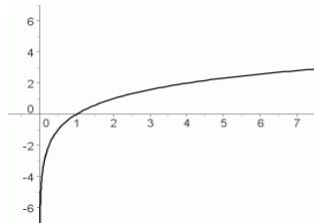
[1]

Logistic Regression – Optimization

- Assume training examples are sampled I.I.D. (Independent and Identically Distributed)
- Likelihood Function: $L(\boldsymbol{\beta}) = \prod_{i=1}^N p(y^{(i)} | \mathbf{x}^{(i)}; \boldsymbol{\beta})$
- Define the loss function as the negative log of the likelihood function, and use gradient descent to solve for optimal $\boldsymbol{\beta}$
 - Negative to make it a minimization problem
 - Log for numerical reasons
- Tune learning rate and regularization parameters


Logistic Regression – Optimization

- $p(y^{(i)}|\mathbf{x}^{(i)}; \boldsymbol{\beta}) = \left(1 - p(\hat{y} = 0|\mathbf{x}^{(i)}; \boldsymbol{\beta})\right)^{y^{(i)}} \left(p(\hat{y} = 0|\mathbf{x}^{(i)}; \boldsymbol{\beta})\right)^{1-y^{(i)}}$
- $L(\boldsymbol{\beta}) = \prod_{i=1}^N \left(1 - p(\hat{y} = 0|\mathbf{x}^{(i)}; \boldsymbol{\beta})\right)^{y^{(i)}} \left(p(\hat{y} = 0|\mathbf{x}^{(i)}; \boldsymbol{\beta})\right)^{1-y^{(i)}}$
- $loss(\boldsymbol{\beta}) = -\log(L(\boldsymbol{\beta}))$
- $loss(\boldsymbol{\beta}) = -\sum_{i=1}^N (y^{(i)}) \log\left(1 - p(\hat{y} = 0|\mathbf{x}^{(i)}; \boldsymbol{\beta})\right) - \sum_{i=1}^N (1 - y^{(i)}) \log\left(p(\hat{y} = 0|\mathbf{x}^{(i)}; \boldsymbol{\beta})\right)$
- $loss(\boldsymbol{\beta}) = \sum_{i=1}^N \log(1 + \exp(-z^{(i)})) + \sum_{i=1}^N y^{(i)} z^{(i)}$
 - Where $z = \boldsymbol{\beta}^T \mathbf{x} + \beta_0$
- Convex function in $\boldsymbol{\beta}$ – therefore we should be able to find the global optimum
 - $\boldsymbol{\beta} \leftarrow \boldsymbol{\beta} - \lambda \frac{\partial loss}{\partial \boldsymbol{\beta}}$



Logistic Regression – Regularization

- Define priors over the weights
 - $\max\{\log[p(\boldsymbol{\beta}) \prod_{i=1}^N p(y^{(i)}|\mathbf{x}^{(i)}; \boldsymbol{\beta})]\}$ w.r.t. $\boldsymbol{\beta}$, where $p(\boldsymbol{\beta}) = N(0, \alpha^{-1}\mathbf{I})$
- Prior biases the weights towards zero -> prevents weights from growing too large




Classification – Other Algorithms

- Popular ones
 - Decision Trees
 - Random Forests -> Xbox Kinect!
 - k-Nearest Neighbour (kNN)
 - Recommender Systems
 - Naïve Bayes
 - Spam Detection
 - Neural Networks
 - Image & Speech Recognition, Machine Translation
 - Support Vector Machines (SVM)
 - Similar use cases as neural networks
 - Mixture of Experts
 - Combines decisions of different algorithms
 - Netflix recommender system!

Validating Results

Cross Validation – Hyperparameter Tuning

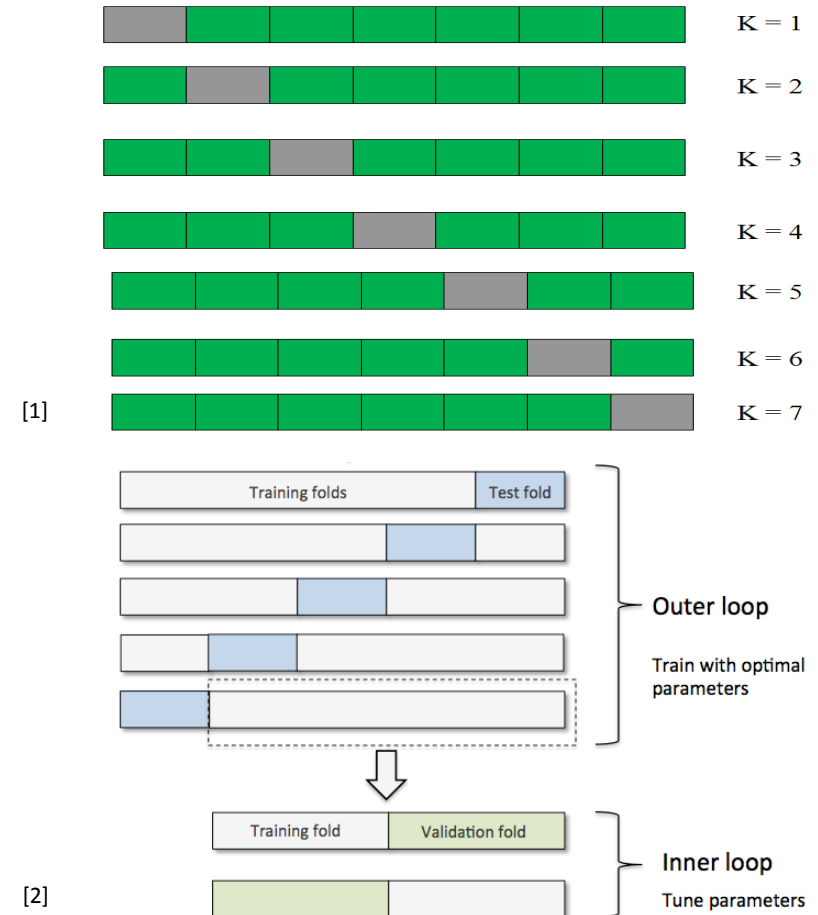
- Need 3 things
 - Training set
 - Validation set
 - Test set



- Use training set to optimize β
- Use validation set to optimize for hyperparameters (ex: λ, α)
- Use test set to evaluate results on unseen data set
 - DO NOT over-test on the test set

Cross Validation - Techniques

- Widely used: k-fold cross validation
 - Increases data set utility
 - Only used to evaluate a single model's performance, not to tune hyperparameters... (see below)
- Alternative: nested k-fold cross validation



[1] <https://www.analyticsvidhya.com/blog/2016/02/7-important-model-evaluation-error-metrics/>

[2] <https://sebastianraschka.com/faq/docs/evaluate-a-model.html>

Metrics for Binary Classification

- Sensitivity: proportion of positives correctly identified
- Specificity: proportion of negatives correctly identified
- $Sensitivity = Recall = \frac{TP}{TP+FN} = \frac{TP}{\text{all ground truth positive cases}}$
- $Specificity = \frac{TN}{TN+FP} = \frac{TN}{\text{all ground truth negative cases}}$
- $Accuracy = \frac{TP+TN}{TP+TN+FP+FN} = \frac{\text{all correct predictions}}{\text{all predictions made}}$

Metrics for Binary Classification

- Recall: fraction of true events that were detected
- Precision: fraction of detections that were true events

- $Recall = \frac{TP}{TP+FN} = \frac{TP}{\text{all ground truth positive cases}}$

- $Precision = \frac{TP}{TP+FP} = \frac{TP}{\text{all cases predicted as positive}}$

- $F1 = 2 \frac{P * R}{P + R}$

Conclusions

- We examined two supervised learning algorithms
 - Linear Regression
 - Logistic Regression
- We examined how to build ML algorithms from the ground up
- We saw that in supervised learning, we use ML algorithms to find β
 - It's still up to us to acquire input data X and ground truth labels y
- We examined how to validate ML models and how to measure performance